

Mobile Information Systems

Lecture 12: Mobile Algorithms

© 2015-23 Dr. Florian Echtler
Bauhaus-Universität Weimar
Aalborg University

Mobile algorithms

- Moving object databases
 - Queries for moving objects
 - Modelling of dynamic attributes
- Synchronization (Dropbox & Co.)
 - CAP theorem
 - Rsync algorithm
 - Trickle algorithm


Moving object queries

- Which restaurant is reachable within 30 min.?
- When will I arrive at the Marktplatz?
- How many ambulances are placed within a radius of 10 km?
- Show a warning if the distance between two airplanes will be less than 500 meters in the next 10 minutes!
- Which UPS driver will be closest to Bauhausstraße 11 at 17:10?

Moving object queries (2)

Image source (FU): © 2015 GeoBasis-DE/BKG (© 2009) Google

- Performing a query may or may not depend on the location of the querying device
- Query is related to the current, a past or a future location of an object
- Space and time might be independent
- Location information has vary-uncertainty
- Objects are (mostly) assumed dimensions)



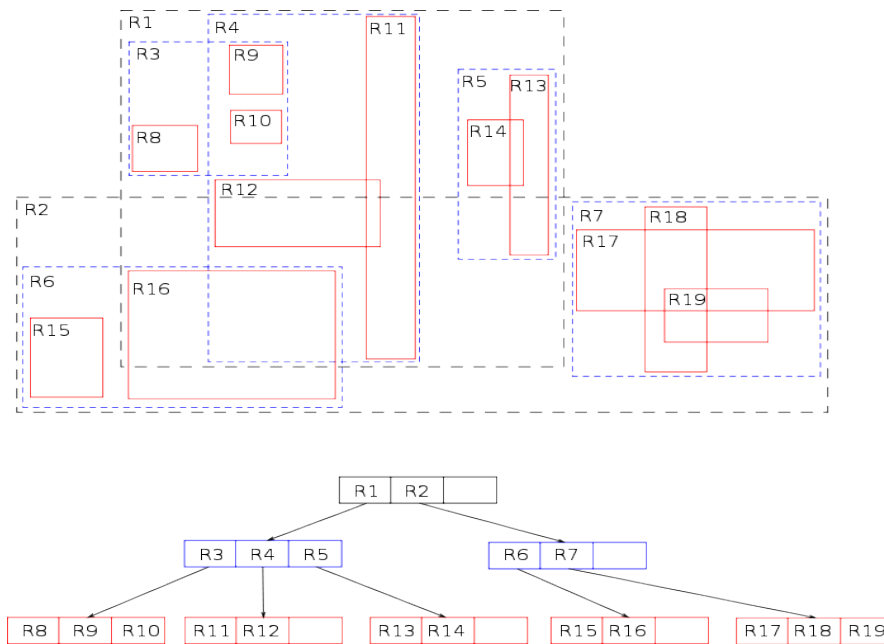
Issues with “traditional” DBMS

- Value is valid until it is changed
- Movement requires permanent updating
 - “Expensive” updates
 - High network traffic
 - Uncertainty
- Typically no history
- No possibility to forecast future locations

Issues regarding indices

Image source (PD): <https://en.wikipedia.org/wiki/R-tree#/media/File:R-tree.svg>

- Possibly large number of moving objects
- Efficient access required
 - Index over location information
- Spatial index: R-tree
 - Each node is a rectangle covering all nodes beneath
 - Updates expensive → ill suited to moving objects



Issues regarding query languages

- Queries with spatial and temporal conditions:
 - Spatial-temporal range query: “objects that overlap with polygon P within next 3 minutes”
 - Spatial-temporal join query: “find airplanes that will have a distance less than 1000 m and return the regarding time point”
- Traditional query languages like SQL are not (well) suited
 - Extensions for spatial/temporal aspects exist, but not optimized for moving objects

Issues regarding uncertainty

- Implications for queries:
 - 2 different results
 - Objects that certainly fulfil the condition
 - Objects that possibly fulfil the condition
 - Or per object: possibility that a condition is fulfilled
- Questions:
 - When does uncertainty become more expensive than updates?
 - How can a DBMS „tell“ how uncertain a piece of information can be at most?

Location model for moving objects

- Traditional:
 - Stored value is valid until it is explicitly changed
- MOD (moving object databases):
 - Explicit updates unacceptable
 - Object might be temporarily disconnected
- Solution
 - Location (and other attributes) are represented as function over time

Dynamic attributes

- Dynamic attribute A consists of 3 sub-attribs:
 - updateValue, updateTime, function $f(t)$
- Value v of A:
 - $t = \text{updateTime}$: $v = \text{updateValue}$
 - $t = \text{updateTime} + t_0$: $v = \text{updateValue} + f(t_0)$
- Explicit update of updateValue and/or $f(t)$:
 - updateTime = timestamp of last change

Dynamic attributes (2)

- Queries for same attribute at t_1 and t_2 may differ, even without intermediate updates
- Queries about past possible
 - Needs log of prior updates
 - Storage space?
- Queries about future possible
 - Prediction based on current function f
 - Changes to f will change query result (e.g. airplanes: deviation from flight plan?)

Location attributes

- Location of a moving object:
 - 2 (or 3) dynamic attributes L.x, L.y, (L.z)
 - $f(t) = \text{updateValue} + (t - \text{updateTime}) * \text{velocity}$
- Suitable for objects with straight movement (airplanes, ships at sea)
 - Many updates for objects that follow a route (e.g. street) and frequently change their direction
 - Each turn requires an update of the function sub-attribute

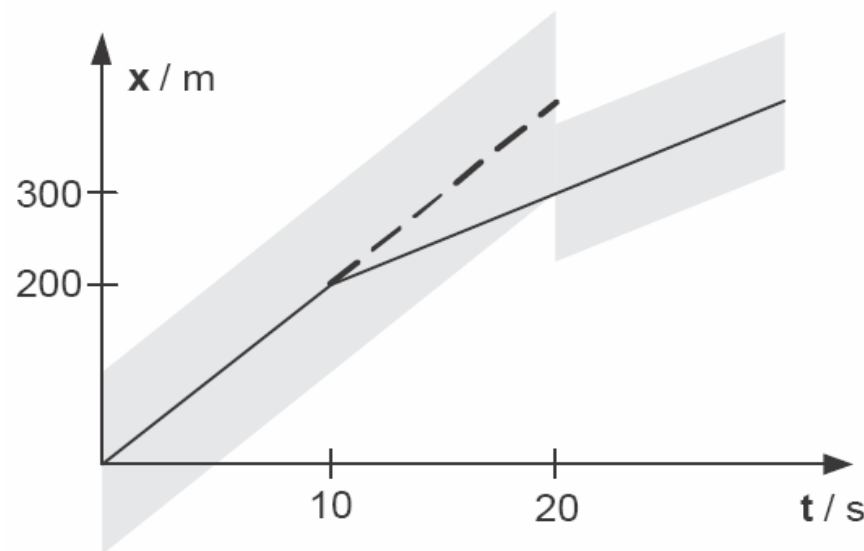
Location attributes (2)

- Better handling of road-based routes
- Location attribute with 4 sub-attributes:
 - updateValue – current index on route
 - updateTime – time point of last update
 - route – list of route points
 - speed – scalar speed
- $f(t) = \text{route}[\text{updateValue} + (t - \text{updateTime}) * \text{speed}]$
 - Requires interpolation for long straight segments

Uncertainty and inaccuracy

Source (FU): Mobile Datenbanken, Höpfner et al., dpunkt.verlag

- Position information is inaccurate
- Objects do not move exactly along the projected route but update only if ...
 - Inaccuracy is above threshold
 - Time interval expires
- From the view point of a query: correctness of result is uncertain



Spatial Indices

- Frequent queries for dynamic attributes
 - Range queries → index on attributes required
- „Normal“ spatial index is not applicable
 - High update frequency
- Division of the problem in two sub-problems
 - Suitable geometric representation of dynamic attributes
 - Time as an additional space dimension
 - Efficient access via index structures
 - Classical spatial indexes

Synchronization

Several issues for Dropbox, OwnCloud & Co.:

- Problem: intermittent connectivity
 - Transfers may be interrupted
 - Updates may happen while offline
- Problem: bandwidth on mobile (still) limited
 - Impossible to just copy GBs of data
 - Adaptive synchronization required

Background: CAP theorem

Source (FU): <http://www.infoq.com/articles/cap-twelve-years-later-how-the-rules-have-changed>

“The CAP theorem states that any networked shared-data system can have at most two of three desirable properties:

- *consistency* (C) equivalent to having a single up-to-date copy of the data;
- *high availability* (A) of that data (for updates);
- tolerance to network *partitions* (P).”

Note: network partitions = disjoint, unconnected sub-networks (single, disconnected device is its own network partition!)


CAP theorem Dropbox?

- CAP theorem → 3 options: CA, AP, PC
 - Which combination applies to Dropbox etc.? (AP)
 - How is the missing feature handled?
- Examples for other combinations?
 - CA – Example: “classic” cluster database (MySQL, MariaDB, PostgreSQL, Oracle, ...) → whole DB is offline when connections within cluster fail
 - CP – Examples: some distributed databases (not as widely used) → e.g. disable “minority” partitions

Synchronization: rsync algorithm

Source (FU): https://rsync.samba.org/tech_report/node2.html

- rsync = Unix tool, created 1996 by A. Tridgell
- Goal: synchronize two *similar* files
 - Designed for slow communication links
 - Core idea: split file into fixed-size blocks, calculate checksum for each block
 - Send checksums to synchronization partner
 - Decide which blocks to transmit




Rsync algorithm (2)

Source (FU): https://rsync.samba.org/tech_report/node3.html

- Works fine if blocks stay in place
- Problem: what if offsets differ?

Node A



Node B




- Inserted data (red) → all following c_i change
- Solution: *rolling* checksum

Rsync algorithm (3)

Source (FU): https://rsync.samba.org/tech_report/node3.html

- Rolling checksum: can be computed for *every possible block offset* in one pass
 - “Sliding window” approach (cf. signal processing)
 - Blocksize s , values v , checksum c , offset i
 - Simplified: $c(i) = (c(i-1) - v[i-1] + v[i+s]) \bmod M$



Rsync algorithm (4)

Source (FU): https://rsync.samba.org/tech_report/node3.html


- During calculation of rolling checksum:
 - Search checksum list received from sync partner
 - Note: list only contains only 1 checksum per block
 - Fast search using hash table possible
 - If match found: verify with *strong checksum* (MD5)
 - If match verified: current block exists on both sides
 - If no match: block needs to be transmitted
- Rsync algorithm is widely used (Dropbox!)
- Also suitable for backups etc. → space-saving

Synchronization: issues?

- Rsync algorithm requires stable connection
 - What if device temporarily offline?
 - What if 2 nodes make changes while offline?

→ (Merge) conflicts

- Resolved using 3-way merge algorithm
- Problem: only applicable to unstructured files, e.g. source code, plain text



Mesh networks: Trickle algorithm


Source (BSD): <https://tools.ietf.org/html/rfc6206>

- Goal: distribute shared state across mesh
 - Examples: sensor configuration, routing tables, software updates, ...
 - Avoid “broadcast storms”, keep packet rate low
- Basic concept for each node: broadcast state, unless other transmissions suggest consistency
 - A broadcasts version V
 - B already has version $V+1$
 - B knows A needs update
 - B broadcasts version $V+1$
 - A has new data

Mesh networks: Trickle algorithm

Source (BSD): <https://tools.ietf.org/html/rfc6206>

- Interval $[I_{\min}, I_{\max}]$, redundancy constant k
- Set random initial timer interval $I := [I_{\min}, I_{\max}]$
- Start timer (interval I), set $c := 0$, $t := [I/2, I]$
 - Receiving "consistent" transmission $\rightarrow c++$
 - Receiving "inconsistent" TX $\rightarrow I = I_{\min}$, restart timer
 - Interval time = t && $c < k \rightarrow$ transmit own state
 - Interval time = I (timer expired) $\rightarrow I = \min(2*I, I_{\max})$



The End

